

# CI583: Data Structures and Operating Systems


## OS Design Principles

# Outline

## 1 Processes and threads

# Processes and threads

Conceptually, the **thread** is a unit of computation, to be stopped and started by the OS.



# Processes and threads


Each thread belongs to an enclosing [process](#), which can also be stopped and started by the OS and which is used to store context which includes:

- an [address space](#) (the set of memory locations that contains the code and data for the program),
- a list of [references to open files](#), and
- other information might be common to several threads.

# Processes and threads

To represent a process we use a data structure called a **process control block** (PCB), containing references to all the info given above and to a list of threads.

Each thread is represented by a **thread control block** (TCB). The TCB contains references to thread-specific context, including the thread's stack and contents of registers.



# Processes and threads

The OS needs to be able to **create**, **delete** and **synchronise** processes.

In Windows and \*NIX a process is either **active** or **terminated**.

Terminated processes are deleted (by a process dedicated to this purpose) when all of its threads are terminated and there are no more references to the process from elsewhere.

# Processes and threads

When a process is created, the address space is loaded into memory. On Linux, this is accomplished by the `exec` program.

How does the OS initialise the address space?

One approach would be to copy the code and data of the program into the address space, but this would be inefficient.

# Processes and threads

The code section of the program is read-only and so can be shared by any processes executing the same program.

The parts of the data section that are never modified can also be shared.



# Processes and threads

A better approach is to **map** the executable file into the address space.

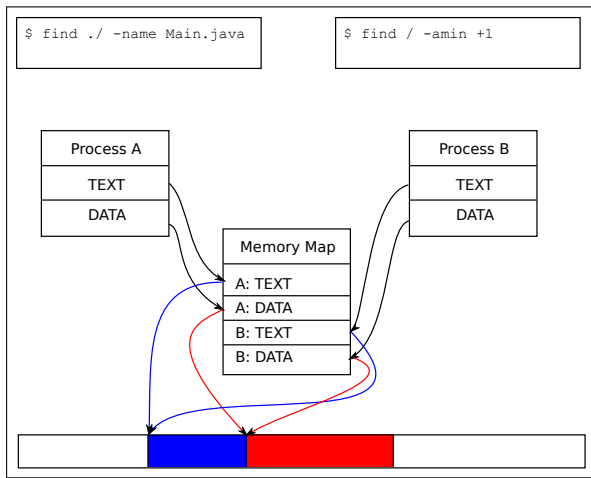
Both the address space and the executable file are divided into blocks of equal size called **pages**.

# Processes and threads

The text regions of all processes running this executable are set up using hardware-address translation facilities, with each process mapping to the same location.

The data regions of each process initially refers to a single copy of the data portion of the executable, which has been copied into memory.

# Processes and Threads



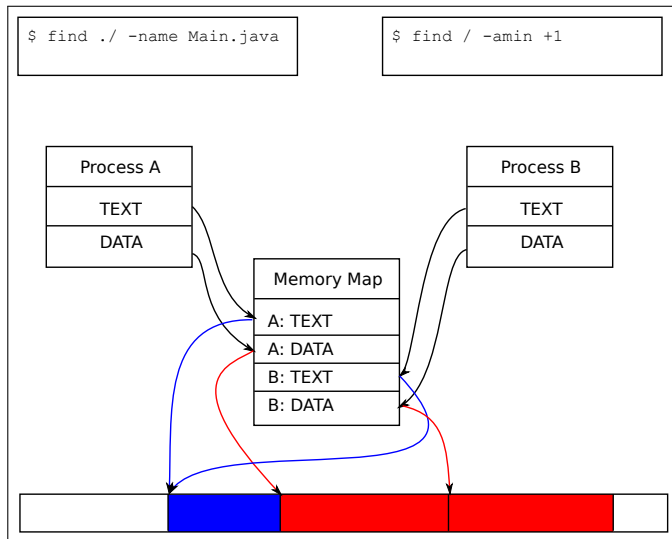
# Processes and threads

When a process modifies data for the first time, it is given a new, private page containing a copy of the pristine page.

This is called a **private mapping**.

Modern systems also use the notion of a **shared mapping**: when data is modified the original page is altered and all processes see the change.

# Processes and threads



# Processes and threads

When a thread is created it is in the **runnable** state.

At some point it is switched into the **running** state by the **scheduler** (which we will come back to in more detail).

A thread that is running may then be put into the **waiting** state, for instance because it has initiated an I/O action and needs to wait for the result before doing anything else.

# Processes and threads

The thread can put itself into the waiting state, or can be moved into the state by the OS.

If the OS uses [time-slicing](#), whereby threads are run for a maximum period of time before relinquishing the processor, the scheduler can move the thread into the waiting state.

# Processes and threads

A thread can move itself into the **terminated** state, or it can be moved there by the OS.

Note that a thread must be in the running state at least once before terminating.

Terminated threads are removed in various ways, such as by a dedicated “reaper” process.



# Processes and threads

