

# CI583: Data Structures and Operating Systems

## OS Design Principles

# Outline

1 Managing hardware

2 Shared libraries

# Managing hardware

In an earlier lecture we discussed device drivers, which contain the code that knows how to interact with a given device.

How do we make sure that the right driver is associated with the right device, that all devices are properly initialised when the system starts, or that new devices are recognised correctly when attached?

# Managing hardware

In early UNIX systems the kernel had hard-coded support for the relevant devices. In order to add a new device, the user needed to recompile the kernel image.

Each device was identified by a **device number** which was a pair of the **major device number**, which identified the driver, and **minor device number**, which identified the particular device among those handled by this driver.

# Managing hardware

Special files were created on the file system in the `/dev` directory to refer to devices.

```

jb259@nowa424-jblin:~/Dropbox/CI283/lectures/lecture3$ ls /dev/
audio          loop4          ram4           tty10          tty39          tty50
block/         loop5          ram5           tty11          tty4           tty51
bsp/           loop6          ram6           tty12          tty40          tty52
bus/           loop7          ram7           tty13          tty41          tty53
cdrom@         mapper/        ram8           tty14          tty42          urandom
cdrw@          mcelog         ram9           tty15          tty43          usbmon0
char/          mem            random         tty16          tty44          usbmon1
console        mixer          rfkill         tty17          tty45          usbmon2
core@          net/           root@          tty18          tty46          usbmon3
cpu_dma_latency network_latency rtc0           tty19          tty47          usbmon4
disk/          network_throughput rtc0           tty2           tty48          usbmon5
dsp            null           sdd0           tty20          tty49          vcs
dvd@           nvdiad0        sda            tty21          tty5          vcs1
dvdrw@         nvdiactl       sda1           tty22          tty50          vcs2
ecryptfs       oldmem         sda2           tty23          tty51          vcs3
fb0            pktdvd/        sda3           tty24          tty52          vcs4
fd@            port           sequencer      tty25          tty53          vcs5
full           ppp            sequencer2     tty26          tty54          vcs6
fuse           psaux         sg0            tty27          tty55          vcs7
hidraw0        ptmx          spl            tty28          tty56          vcsa
hidraw1        pts/          /              tty29          tty57          vcsa1
hidraw2        ram0          snapshot       tty3           tty58          vcsa2
hidraw3        ram1          snd/           tty30          tty59          vcsa3
hpet           ram10         sndstat@       tty31          tty6           vcsa4
input/         ram11         sr0            tty32          tty60          vcsa5
kmsg           ram12         stderr@        tty33          tty61          vcsa6
log=           ram13         stdio@         tty34          tty62          vcsa7
loop0          ram14         stdout@        tty35          tty63          vga_arbiter
loop1          ram15         tty            tty36          tty7           zero
loop2          ram2          tty0           tty37          tty8
loop3          ram3          tty1           tty38          tty9
  
```

These “special” files refer to the device using its device number. Thus, when an application opens the “file” `/dev/sda1` the kernel uses the right device and driver.

# Managing hardware

This approach was laborious and inadequate for the number of devices that can be used with today's systems.

The modern approach uses what we might call [meta-drivers](#), each responsible for the class of devices that can use a particular bus,

So, the USB driver knows about USB devices and probes the system at start-up time to see which are attached.

# Managing hardware

The correct drivers and kernel modules are then loaded to initialise these devices.

Whilst a modern Linux system is running, the `udev` “daemon” listen for the connection and removal of devices and updates the contents of `/dev` accordingly.

Most of what `udev` does is in userland (doesn't require special privileges).

# Interacting with a device

We begin by considering how the OS might interact with a very simple device – a **terminal**, which takes user input at the command line and displays the results in a simple text-only UI.

Although this might seem like an obsolete example, the way that the OS reads from and writes to this device is something that can be adapted for many contexts.





# Interacting with a device

We can see straight away that it's not enough to read one character at a time from the device, or to write data straight to the device. Why not?

- ① Data may be sent to the device faster than they can be processed, or generated by the user faster than the application can accept it.

# Interacting with a device

- ② Chars may arrive from the keyboard when there is no waiting read request.

## Interacting with a device

- ③ Input may need to be processed in some way before they read the application.

For instance, chars need to be echoed to the screen so that the user can see what they are typing.

Chars may be grouped into lines which can be edited before submitting by pressing enter.

The terminal may support tab-completion.

# Interacting with a device

Items 1 and 2 are examples of the **producer-consumer** problem, in which two processes need to synchronise their access to a finite **queue** or **buffer**.

The **producer's** job is to generate a piece of data, put it into the buffer and start again.

# Interacting with a device

At the same time, the **consumer** is consuming the data (i.e., removing it from the buffer) one piece at a time.

The problem is to make sure that the producer won't try to add data into the buffer if it's full and that the consumer won't try to remove data from an empty buffer.

## Interacting with a device

The OS will communicate with many devices using the same model of input and output queues (even a window manager, which controls the gui – here the input comes from keyboard and mouse, and the window manager needs to keep track of which application window has the focus, whilst the output is a stream of requests to repaint parts of the display using new data).

We separate this common functionality into something called a **line discipline** module.

# Shared libraries

We have already mentioned the fact that each OS has an API and that high-level languages provide convenient calls to the API in [standard libraries](#).

As well as wrappers for the API, standard libraries contain large numbers of standard functions for convenience.

# Shared libraries

These shared libraries are called **dynamic-linked libraries** on Windows and **shared objects** on Linux.

One advantage of using them is that they need not be loaded until needed, improving the start-up time of a program.



# Shared libraries

The biggest advantage is **code reuse** – few programmers would attempt GUI programming without a library.

One disadvantage is that they bring increased complexity (DLL A depends on DLL B, which depends on DLL C – are the versions of these DLLs compatible?) – this complexity is greatly tamed by managed runtime environments like .NET or the JVM.

# Shared libraries

When a program is converted into machine code that can be executed directly by the OS, it needs to contain everything it needs to run.

When we include a call to a shared function in our program, such as a call to the C function `printf` (or, more indirectly, `System.out.println` in Java), a compiler could take one of several choices:

# Shared libraries

Approaches to shared libraries:

- ❶ Include a copy of `printf` with our program. This would make our programs much larger than they need to be and be a waste of storage.
- ❷ Load a copy of `printf` at runtime, along with our program. There would only be one copy on disk but each program would load it into memory: again, this would be very inefficient.

# Shared libraries

Approaches to shared libraries:

- ③ Ensure that all processes that make use of `printf` refer to the same known location,  $L$ . If some process wants to make use of that location for another purpose, move the single copy of `printf` to another location,  $M$ . This is the approach taken by Windows.
- ④ Ensure that each process that calls on `printf` maintains a table mapping an arbitrary memory location,  $L_1$ , to the real location in memory of `printf`,  $L_2$ . This approach is called [position-independent code](#) and is the approach used by most modern UNIX systems.

# After the break

Storage, virtual memory, virtualisation, scheduling.